

# Are Lexical Representations Separable? Item and Category in Learned Distributed Representations

Anonymous ACL submission

## Abstract

PLACEHOLDER — author to write.

## 1 Introduction

## 2 Experiment 1: What a Learner Builds from a Known Language

The goal of experiment 1 is to validate our measurement on an artificial language whose structure we control as well as examine how a learner represents that structure.

### 2.1 Dataset

In order to examine the representation of linguistic concepts, we constructed an artificial language which comprises a probability distribution over a vocabulary of 600 tokens. Every observation is a combination of a **form** (one of 60), a **class** (one of 4), and a **context** (drawn from a pool of 400). The motivation for this setup is to give each word a naturalistic environment: the word occurs in a context that favors a specific class-membership (e.g., *walk* as a verb occurs in a different context than *walk* as a noun). Each form is observed 24 times in each class, every time in a different context, which is what a word in a corpus looks like: one type, many occurrences, each in different company. Consider *walk* used as a verb, in one of those contexts. Four things determine the distribution over the next token. *Walk* itself contributes the same amount whichever class it is used as, and the verb class also contributes the same amount for every verb. The third contribution is comprised of the parts of the representation specific to *walk*-as-a-verb over and above those two: every form and every class has its own random vector, and this term multiplies *walk*'s by the verb's, entry by entry. A sum could be pulled back apart into a form part and a class part; a product cannot, which is what makes this term genuinely joint. The context contributes the

rest, which is noise as far as our question goes, though the model still has to encode it exactly as a language model has to encode the rest of a sentence. The four are weighted, summed and passed through a softmax:

$$P(\text{next token} \mid f, c, x) = \text{softmax} \left( s \left[ w_{\alpha} \text{item}(f) + w_{\beta} \text{class}(c) + w_{\gamma} \text{int}(f, c) + w_{\text{ctx}} \text{context}(x) \right] \right) \quad (1)$$

Each term is scaled to unit length before it is weighted, so changing the weights changes the mix of the three effects without changing how much signal there is in total, and the task stays equally hard everywhere in the grid.

The weights  $w_{\alpha}$ ,  $w_{\beta}$  and  $w_{\gamma}$  in Equation 1 control how much of the true distribution each of the three effects contribute, and we cross them factorially over  $\{0, 0.5, 1, 2\}$ : zero removes an effect from the language altogether, and the other three values make an effect half as strong as another, equal to it, or twice it. Similarly,  $w_{\text{ctx}}$  governs how much the context contributes, so we cross those languages with three context weights, ranging from context being a minor part of the true distribution to the dominant one. We train a model on each of the resulting languages at hidden widths of 8, 16, 32, 128, with either a linear or a ReLU activation (Section 2.2). The whole grid is then run with five random seeds, each redrawing the language and reinitializing the model, giving 7,560 runs in all, of which 99.7% trained to a loss plateau.

It is worth noting that entanglement is guaranteed under a specific subset of the languages. That is, the structure needs 16 dimensions, so its ratio to the hidden width, which we call **capacity**, runs from 0.125 to 2.0. Below one there is room to give every component its own axis, so any entanglement was chosen; above one there are more components

than dimensions to put them in, so they cannot all have their own axis and entanglement is forced (Elhage et al., 2022).

## 2.2 Method

All code and analyses in this paper can be found here: <https://github.com/znhoughton/representation-separability>.

### 2.2.1 The learner

The input to the learner model comprises three pieces of information: the form, the class, and the context. Each has its own embedding, and the three are concatenated and passed to a standard multi-layer perceptron (MLP) with a single hidden layer of width  $d$  and either a linear or a ReLU activation function, which outputs a distribution over the vocabulary (Figure 1). The training objective is thus to minimize cross-entropy between the model’s predicted distribution over the vocabulary and the true distribution.

### 2.2.2 Decomposing a representation

We measure separability by decomposing a word-in-context representation into its lexical representation, its class representation, and its idiosyncratic representation. The representation is the MLP’s hidden layer, which for a given form and class varies across the contexts it is observed in, so let  $M[i, c] \in \mathbb{R}^d$  be the mean of those hidden states for item  $i$  used at class  $c$ . We decompose that grid by centering, so that

$$M[i, c] = \mu + \alpha(i) + \beta(c) + \gamma(i, c), \quad (2)$$

where  $\mu$  is the grand mean;  $\alpha(i)$ , item  $i$ ’s mean across classes minus  $\mu$ , carries the lexical information;  $\beta(c)$ , the mean across items at class  $c$  minus  $\mu$ , carries the class information; and  $\gamma(i, c)$  carries the idiosyncratic information, what item  $i$  at class  $c$  holds that neither marginal accounts for (Figure 2). Because every item is observed in every class with equal weight, neither marginal absorbs part of the other. The four terms add back to  $M[i, c]$  exactly, and no two of them vary together across the grid, so their squared lengths sum to the total and we report each as a share of it.

Each of these averages is taken over finitely many observations, which differ only in context, so each carries some of the context term of Equation 1 as noise. We remove it by dividing the observations

behind each average into two even-sized splits, running the decomposition on each split separately, and estimating each component’s size as the dot product of its two half-estimates, which is large when the splits agree about the component and near zero when they do not. It is the contexts that are divided and not the words, so every word appears in both splits, and the dot product compares two estimates of the same word built from different sentences: noise that differs between the splits is independent and cancels.

For each representation we report two measurements. The first measures the fraction of the representation accounted for by lexical, class and idiosyncratic information: each component’s squared length, summed over its vectors and divided by the total, so that the three fractions sum to one. The second measures whether those components occupy the same directions. In order to do so, we project the item vectors onto the class direction and record the fraction of their squared length that survives, and we project the interaction onto the smallest subspace containing both the item and the class vectors (i.e., the span of the two subspaces). A representation is separable if a word’s representation is the sum of an item part and a class part with the two on their own axes, so it fails either when the components are non-orthogonal, so that projecting out the class also removes item information, or when there is a third component that is neither an item part nor a class part, whatever its orientation. The question is accordingly not whether class information is present, but how it is carried: on a single direction shared across words, or on a separate direction for each of them.

To quantify whether the effect of size is meaningful, we test each against zero. Size is estimated by taking the dot product of two independent half-estimates of the same component, so noise that differs between the halves cancels out and its value is zero in expectation when the component is absent, which is what makes zero the null. We repeat the measurement over two hundred re-samplings to estimate how much the size varies with which observations fall in which half, and we call a size meaningful when its confidence interval excludes zero. On representations with a component planted at exactly zero, this flags an effect about 5% of the time (Appendix A), the false-positive rate expected of a 95% confidence interval.

In order to quantify whether the overlap is meaningful, we measure the overlap an effect has with a



Figure 1: What the learner sees. Every form is observed 24 times at each class, each time in a different context; that variation across contexts is what the measurement later splits in two.

random subspace. We do so by holding the effect exactly as measured and projecting it onto a random subspace of the same rank, two hundred times, which leaves the effect untouched and varies only the subspace. These two hundred projections give us a distribution of overlaps that an unrelated subspace would yield. An effect lower than this baseline indicates meaningful orthogonality; an effect larger than it indicates meaningful non-orthogonality. We present each overlap as a multiple of the median of this distribution, together with how many measurements fall outside its 95% interval. This gives both the size of the effect and how consistently it holds across measurements.




Figure 2: A  $2 \times 2$  grid of item-class means. Under an item part plus a class part, three corners fix the fourth and the parallelogram closes. The arrow is  $\gamma(i, c)$ , the displacement between the fourth mean and the point that additive account places it at, with length  $\|\gamma(i, c)\|$ .

### 2.3 Results

The weights in Equation 1 describe the true distribution the model is trained to match, while the measure reads the representation it builds, so their

numbers do not match: a learner handed a language that is 85% interaction can build a representation in which the interaction is a far smaller share, because the representation only has to be enough to produce the answer. Validation asks that the measure track what was planted, read zero when zero was planted, and have a known floor. Appendix A reports the separate check in which the components are planted directly in a representation.

The linear condition reads zero only where the language is additive. Across all 2,880 runs where it contained an interaction, the linear model's interaction interval excludes zero in 36% of them, rising from 8% at the weakest planting to 94% at the strongest, and 2% where none was planted. It may seem contradictory for a linear model to carry an interaction effect; it is contradictory only for what the model can compute, not for what its representations carry. The reason is that the learner's representations are not context-invariant. Which contexts *walk* occurs in is partly a fact about *walk*, partly about noun-hood, and partly about the two together, and the part that varies by word and class jointly cannot be split back into an item part and a class part. The form and class vectors are shared across every pairing, so in a linear model the contexts a form is observed with are the only place a property of the form-as-class can live; its hidden layer is a sum of its inputs, so it preserves whatever the contexts do differently in the two classes, carrying the property without ever representing it. What ReLU adds is another place to put it: a hidden unit responds to the word and the class together, so the model can compute the property from them rather than take it in from context. The ReLU condition flags 77% of the runs where the language contained an interaction against 36% for the linear, and since the linear model's own rate rises with the planting, the gap between the two is widest in the middle of the range and closes almost entirely at the top.




Figure 3: How the item-by-class grid divides in the learner’s representation, by hidden width, under ReLU. The three shares sum to one within each run. Each panel holds the runs whose language carried a given amount of interaction. Points are medians over runs; bars are the 2.5th and 97.5th percentiles across runs, so they show how differently two runs of the same language at the same width come out. Runs in which none of the three components had a re-split interval excluding zero are not plotted, since the three shares then divide noise; all such runs are at width 8, and their number is given in the text.

Because the measure recovers components planted directly into a representation almost exactly (Appendix A), what follows can be read as a fact about the learner rather than about the instrument. The question is not whether the measure finds the interaction that was planted. It is what a learner does with a language whose composition we know.

What it does depends on how much room it has. Pooled over languages, the interaction takes 0.002 of the grid at a hidden width of 8 and 0.068 at 128, a 45-fold rise, while the item and class effects barely move above width 16 (Figure 3). The interaction is the component that grows into the space a wider representation provides, and the two main effects are what a narrow one spends its dimensions on first. The rise is steepest where there is most to build: in languages that are mostly interaction it runs from 0.028 at width 8 to 0.140 at 128, against a far flatter climb where the language holds little interaction to begin with. At the widest representation we train, the interaction’s share is still rising.

That is the reading Experiment 2 tests. If the interaction is gated by capacity rather than by anything about the distinction being represented, then models an order of magnitude wider than these should carry more of it, not less.

Width also decides how reproducible the composition is. At a width of 8 the interval on the item effect spans 1.00 of the available range and the class effect 1.00, so two runs given the same language can divide their representations in almost

opposite ways. By width 32 those intervals have narrowed to 0.63 and 0.61, and they do not narrow further at 128. Narrow representations are not simply poorer versions of wide ones; they are erratic, and which of the item and the class effect dominates is close to a matter of the seed. The same width accounts for every run the figure leaves out: at a width of 8, 23% of runs produce no component whose interval excludes zero, so there is no measured structure for the three shares to divide, while at every wider setting that happens in no run at all.

Item and class effects entangle only when there is no room for them not to, and this is the substantive result (Appendix B). With capacity well below one, ReLU’s item-class overlap is 0.0092 against a chance level of 0.0297, so the two effects are *more* orthogonal than two arbitrary directions would be, which is what a learner that has room and uses it looks like. At capacity 1 the overlap is 0.402 against a floor of 0.2954, about 1.4 times chance, and at capacity 2 it is 0.512; the linear condition stays at or below chance throughout. This is superposition in a setting where we can watch it happen: when the number of directions the structure needs exceeds the number available, the components stop being orthogonal because they cannot all be, and nothing about the language changed across those runs.

The measure also works when the representation primarily encodes contextual information unrelated to the specific form and class. When the context term accounts for more than 80% of the variance

in the logits, leaving under a fifth for the form, the class and their interaction combined, ReLU’s measured interaction, as a share of the grid, holds at 0.033 against 0.038 when context is minor, and the flag rate falls from 96% to 53%. Power degrades, as it must with twenty-four observations per item and class, but the estimate does not drift. Context does compete with the interaction for room, though: at a width of 128 the interaction takes a larger share of the grid in languages that are under 30% context than in those over 70% (Appendix ef{sec-appendix-capacity}). That is the regime Experiment 2 sits in, since a word’s hidden state varies mostly with the sentence it appears in rather than with the word or its class.

Two things carry into Experiment 2. The interaction is the component that grows with room, so a representation an order of magnitude wider than the widest trained here should carry more of it rather than less, which is a prediction the language models can be held to. And whether the item and class effects come out orthogonal is governed by the same capacity rather than by anything about the language, so near-orthogonality in a wide model is uninformative on its own and has to be read against its own chance level.

### 3 Experiment 2: Item and Category in Language Models

Experiment 1 established that a learner fitting an interactive language cannot represent it as item plus class, that even a learner which cannot compute the interaction stores a trace of it in the contexts its forms occur with, and that a nonlinearity additionally makes the components non-orthogonal. Experiment 2 applies the same decomposition to the contextual representations of pretrained language models, across three linguistic distinctions chosen to span grammar and meaning.

#### 3.1 Dataset

For part of speech and grammatical role we concatenate six general-domain English Universal Dependencies treebanks (Nivre et al., 2020): EWT, GUM, LinES, ParTUT, GENTLE and PUD, giving 41,095 sentences and 706,323 tokens. A lemma must be attested in both classes often enough to estimate a mean at each, and a single treebank yields too few such lemmas. For metaphor we use the VUA20 corpus (Leong et al., 2020; Steen et al., 2010), restricted to nouns, verbs, adjectives, and

| distinction      | levels                    | items | tokens | median |
|------------------|---------------------------|-------|--------|--------|
| part of speech   | NOUN / VERB               | 78    | 4,722  | 18     |
| grammatical role | <i>nsubj</i> / <i>obj</i> | 49    | 2,150  | 17     |
| metaphor         | literal / figurative      | 180   | 14,871 | 23     |

Table 1: The stimuli. All three constructions key the item on the **surface form**, so the token is identical in both classes of the distinction. *Items* counts the forms attested at least ten times in each class; *grid tokens* are the tokens those items contribute, and *median* the tokens behind one item-class pairing.

adverbs in VUA’s own annotation, which supplies 88,509 labeled tokens; the adverbs are those of phrasal predicates like *set up*, which VUA annotates as metaphorical in their own right. We measure two model families at three matched scales: OPT-BabyLM, trained on developmentally plausible quantities of child-available language (Warstadt et al., 2023), and Pythia, trained on the Pile (Biderman et al., 2023), at 125M/160M, 350M/410M and 1.3B/1.4B parameters, with widths of 768, 1024 and 2048. Appendix D gives parameter counts and training-token estimates.

#### 3.2 Stimuli

The grids are built from the words themselves, and what they are is worth stating. Each construction keys the item on the **surface form**, lowercased, so every item is same-token: the surface form is identical across the two classes, the representation is read from the same token in both, and no part of any difference can come from tokenization. An item enters a grid only with at least ten tokens at each class, so that a mean can be estimated and then split in two; Table 1 gives the surviving items, the tokens they contribute to the grid, and the median tokens behind each pairing. `scripts/llm/dataset_stats.py` also writes the per-item lists with their per-class token counts to `stimuli_items.csv`, so the stimuli themselves can be inspected word by word.

The part-of-speech grid is the restrictive one: the interaction  $\gamma(i, c)$  is estimable only for an item observed in more than one class, since it is defined by how an item’s representation shifts across classes, and the balance requirement selects forms genuinely used both ways, such as *run* or *hit*, including pairs like *claims* where a noun plural and a verb’s third person are the same token. Role and metaphor carry no such restriction, since any noun occurs as both subject and object and any content

word occurs both literally and figuratively, which is one reason they broaden the claim. The median pairing holds 18 tokens for part of speech, 17 for role and 23 for metaphor, comfortably inside the 10-to-100 observations per pairing that Appendix A measures over.

Keying the item on the *lemma* instead admits more items (130 rather than 78), but a lemma pools every inflected form: *run* and *runs* as a noun against *run*, *runs*, *ran* and *running* as a verb. Those are different tokens, and the consequence shows up at the embedding layer, which has no context to work with and so can only distinguish classes that differ in their input. Keyed on the form the layer-0 interaction is 0.0009 to 0.0035 across models, essentially nothing, which is what a token identical in both classes should give. The difference is inflection rather than representation, and keying on the form removes it at the cost of a smaller item set. Appendix C shows what the measure does on a distinction where the property cannot be arranged at all.

### 3.3 Method

The decomposition, the two reported quantities and the replication logic are those of Section 2.2. For every token we take the hidden state of the word’s **last subword** at every layer, from the embeddings through the final block, aligning subwords to corpus tokens with the tokenizer’s word alignment. Sentences are drawn in a fixed random order and processed until roughly 300K tokens are extracted per model, so all models see the same material, and representations are extracted once with every analysis running offline from them.

The two independent estimates come from the corpus itself: each item in a class is a word used across many sentences, so we divide its token contexts into two disjoint, even-sized splits and estimate the grid from each. What cancels is context noise, the variation in a word’s hidden state across the sentences it appears in, which is the same thing the split cancels in Experiment 1.

Transformer hidden states contain a few *rogue* dimensions whose variance exceeds the rest by two to three orders of magnitude (Timkey and van Schijndel, 2021), and these would otherwise dominate any subspace estimated from the raw scatter, so we divide each dimension by its standard deviation before decomposing. This is a diagonal rescaling rather than a whitening, so it needs no favorable ratio of samples to dimensions; the measure is ac-

cordingly invariant to per-dimension rescaling but not to arbitrary linear reparameterization, which we return to in Appendix E.

### 3.4 Results

The two main effects sit on separate axes. At the deepest layer, item-class overlap runs 0.0032 to 0.0067 for part of speech and 0.0025 to 0.0062 for grammatical role, so projecting out the class direction leaves better than 99% of the item effect intact. Both sit above the upper end of their null, so the residual alignment is real rather than a by-product of the dimension count, but it is small enough that these representations do not achieve class structure by borrowing the directions that carry word identity. Where a shared class direction exists, it is its own direction.

The interaction, however, does not have axes of its own, and this is what decides whether a third component could simply be added to the description and separability recovered in a wider sense. Across the six models, part of speech puts 0.096 to 0.239 of the interaction in the span of the two main effects, role 0.084 to 0.179 and metaphor 0.128 to 0.267, against nulls whose upper ends run 0.023 to 0.058, 0.014 to 0.038 and 0.020 to 0.060 (Table 14). The smallest of these overlaps still exceeds the largest of these nulls, so the pattern rests on every model rather than on their average. The partition of the grid says nothing about direction, so these numbers are not implied by the decomposition: the item-by-class term reuses the very directions that carry item identity and class membership. Read with the first overlap, they say something more specific than “not separable”. The failure is not that item and class have been collapsed onto shared axes, since those remain distinct; it is that the part of class structure specific to each word is written in the same directions as everything else. For metaphor the orientation cannot be read at all, because the shared class direction is too small to fit a subspace to, which is itself the first indication of how little of metaphoricity is carried by a shared direction. The class effect is a cross-split dot product, so its null is exactly zero, and it clears zero in 6 of 6 for metaphor, as it does in 6 of 6 for part of speech and 6 of 6 for grammatical role. All three therefore have a shared class direction that replicates across splits, and what separates them is its size: the class effect is 0.057 to 0.099 of the representation for part of speech and 0.047 to 0.078 for grammatical role, against 0.001 to 0.006 for




Figure 4: The three components through the network, for every model and distinction. Relative depth runs from the static embedding layer (0) to the final block (1). Top: the interaction’s share of representational structure. Middle: the class effect’s share. Bottom: the portable fraction, the class effect divided by the sum of the class effect and the interaction. Each line is one model: color distinguishes the training corpus, linetype the three matched widths. The shaded band is that model’s re-split 95% interval.

metaphor, more than an order of magnitude smaller within every model (Table 14). Metaphor’s shared class direction is real but slight, and most of what distinguishes its uses is carried word by word rather than in a direction shared across them.

A substantial interaction is there, and depth builds it. In all three distinctions the interaction is essentially absent at layer 0, between 0.0004 to 0.0044 across every construction and model, and rises with depth to between 0.0228 to 0.2450 (Figure 4, top row). Because the item is keyed on the surface form, the token is identical in both classes, so the static embedding layer has nothing to distinguish them with, and it does not. Everything the measure finds is therefore created by contextualization rather than inherited from the input, and it is not positional: zeroing OPT-BabyLM’s absolute position embeddings leaves the class structure more word-specific rather than less (Appendix G). This is the component that breaks separability, since it is neither a fact about the word nor a fact about the class.

All three sizes are shares of the grid, and the grid has already averaged away how a word varies across the contexts it appears in, so how much of

the representation the grid is in the first place is worth stating. At the embedding layer it is 0.83 to 0.85 of the variance across the three constructions, which is the same-token design working as intended: identical tokens differ little before anything contextualizes them. At the deepest layer it is 0.32 to 0.33, so most of what the model computes there is about the sentence a word is in rather than about the word and its class. This one quantity separates by family rather than by distinction: OPT-BabyLM’s grid accounts for 0.48 of the representation and Pythia’s for 0.31, in all three constructions, so the Pile-trained models devote more of the deepest layer to context than the developmentally scaled ones do. Multiplying the two through, the deepest-layer item effect for part of speech is 0.25 of the representation rather than the 0.75 it is of the grid, and the class effect and interaction scale down in the same proportion, with context taking the remainder. Both figures are corrected for the noise each grid entry carries as an average over finitely many observations, which otherwise inflates the grid’s apparent share by an amount that shrinks as tokens per item grow and so differs across constructions (Appendix E).

The class effect grows with depth as well (Figure 4, middle row), and that matters for the framing. Models are not choosing between abstraction and entanglement; contextualization builds both at once. A description of these representations as either abstract or exemplar-like is therefore underdetermined by the data, and the useful question is not which one they are but what the balance is.

That balance is graded across distinctions. Table 14 reports it as the *portable fraction*: of the structure that distinguishes the two classes, how much is a direction shared across words rather than specific to each word. Grammatical role is the most portable at 0.40 to 0.58 across the six models, part of speech next at 0.25 to 0.47, and metaphor by far the least at 0.04 to 0.08, an ordering that holds within every model rather than only on average, so a direction shared across words carries barely a tenth of what distinguishes metaphor’s classes. The ordering is worth noting: part of speech has the largest class effect of the three in absolute terms, but also much the largest interaction, so proportionally less of what distinguishes its classes is portable than for grammatical role. Metaphoricity is close to being encoded word by word, and what makes a use of *plant* figurative is largely not what makes a use of *course* figurative. This grading runs across distinctions rather than across models: neither quantity varies systematically with parameter count, and the only consistent family difference is that Pythia builds a larger interaction than OPT-BabyLM at matched width (Figure 4).

## 4 Experiment 3: Does the Interaction Carry the Class?

Experiment 2 shows that the interaction is large and that it replicates across independent splits of a word’s occurrences, but neither shows what it is *about*. An item-by-class term captures any reliable difference between a word’s uses in the two classes, and some of that difference is context correlated with the label rather than the label itself, since metaphorical uses occur in different sentences from literal ones and subjects appear in different constructions from objects. So  $\gamma$  is an upper bound on class-binding, and the question is whether any of it is the class.

### 4.1 Method

We test it by decoding. If  $\gamma$  encodes the class, the class should be recoverable from the part of a

| distinction      | model       | words | class             | null              |
|------------------|-------------|-------|-------------------|-------------------|
| part of speech   | BabyLM 1.3B | 78    | 0.93 [0.91, 0.94] | 0.52 [0.49, 0.55] |
| part of speech   | Pythia 1.4B | 78    | 0.94 [0.93, 0.95] | 0.50 [0.48, 0.53] |
| grammatical role | BabyLM 1.3B | 49    | 0.85 [0.82, 0.87] | 0.50 [0.47, 0.53] |
| grammatical role | Pythia 1.4B | 49    | 0.84 [0.82, 0.87] | 0.52 [0.49, 0.55] |
| metaphor         | BabyLM 1.3B | 180   | 0.65 [0.63, 0.66] | 0.49 [0.47, 0.50] |
| metaphor         | Pythia 1.4B | 180   | 0.70 [0.68, 0.71] | 0.49 [0.48, 0.51] |

Table 2: Class decoded from the interaction, at the deepest layer of the largest model in each family. The item is held fixed and the shared class direction is projected out. *Class* is the accuracy of that decode and *null* repeats it with class labels permuted within each word; brackets are percentile bootstrap intervals over words. Chance is 0.50.

token’s representation that is neither the item effect nor the shared class direction, and we remove both explicitly. All decoding happens *within* a single word, so we ask whether *plant*’s figurative uses can be told from *plant*’s literal ones, never whether *plant* can be told from *account*; the item effect is constant within a word and so cannot contribute. We then estimate  $\beta$  as the difference between the two class means pooled over all items, normalize it, and remove that component from every token, so that whatever separates the classes afterwards is not the direction they share.

What remains is the interaction: the part of a word-at-a-class that is neither a fact about the word nor a fact about the class. We decode it with a cross-validated logistic probe, five folds, classes balanced by subsampling within each word so that chance is 0.50, and average the resulting accuracies over words, with a word contributing only if it has at least eight tokens in each class. Accuracies are averaged unweighted over words and some words rest on few tokens, so the mean carries a percentile bootstrap interval over words. The whole procedure is then repeated with the class labels permuted within each word, which preserves every word’s token count and class balance while destroying the labelling, and that null is where a probe reading its own cross-validation folds would land. We run this at the deepest layer of the largest model in each family, for all three constructions.

### 4.2 Results

The interaction is about the class. In every construction and both families the class is recoverable from the interaction at well above chance while the permuted null sits on chance (Table 2), reaching 0.94 for part of speech, 0.84 for grammatical role and 0.70 for metaphor, with bootstrap intervals ex-

cluding the nulls by a wide margin in all six cases. The null landing on chance is what makes those accuracies interpretable, since it is the same probe, the same folds, the same word-level averaging and the same small per-word samples with only the labelling destroyed.

The effect is graded the same way the interaction’s size is. Part of speech decodes best, role next, metaphor least, matching the order in which the interaction’s share falls across the three in Table 14. The distinction whose interaction is smallest is also the one whose interaction carries the class least strongly, so the geometric and functional measurements agree on the ranking.

Metaphor is the case where this matters most. Its shared class direction is real but slight, too small to fit an orientation to (Section 3.4), so a method reading only that direction would take metaphoricity to be barely represented, yet it decodes from the interaction at 0.70 against a null of 0.49. It is represented; most of what distinguishes its uses lives in the interaction rather than in a direction shared across words, which is exactly what a method looking only for a shared direction would miss.

This result has a definite scope. It shows that  $\gamma$  carries the specific label rather than generic correlated context, since generic context would not survive holding the word fixed and removing the shared axis. It does not show that  $\gamma$  is *only* the label: context correlated with the class within a word survives both removals, so the interaction remains an upper bound on class-binding, and this raises the floor rather than closing the gap.

## 5 Discussion

### Limitations

Everything here is English, and rests on three distinctions measured in two model families at three scales each. The graded ordering we report is an ordering across *these* distinctions, and whether it holds for aspect, animacy or definiteness is untested.

The measure is invariant to per-dimension rescaling and to rotation, but not to arbitrary linear reparameterization, so overlap between the item and class effects is relative to the model’s own basis. We take that to be the basis of interest, since LayerNorm and the feed-forward nonlinearities act coordinate-wise and a downstream output layer operates on the unwhitened space; a reader who wants a fully gauge-free quantity will not find one here.

The component sizes, their growth with depth and whether their intervals exclude zero do not depend on that choice.

Two things we would like to have done. These are endpoint measurements, so they say nothing about the order in which the two components are learned, which needs training checkpoints. And the interaction captures any systematic within-word difference between the classes, so a distinction realized through the contexts a word appears in is measured together with those contexts. For a contextual model we take that to be what representing the distinction consists in, but separating the two would need a construction in which they can be varied independently.

## References

- Stella Biderman, Hailey Schoelkopf, Quentin Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. 2023. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*.
- Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, and 1 others. 2022. Toy models of superposition. *Transformer Circuits Thread*.
- Chee Wee (Ben) Leong, Beata Beigman Klebanov, Chris Hamill, Egon Stemle, Rutuja Ubale, and Xinyang Chen. 2020. A report on the 2020 VUA and TOEFL metaphor detection shared task. In *Proceedings of the Second Workshop on Figurative Language Processing*.
- Joakim Nivre, Marie-Catherine de Marneffe, Filip Ginter, Jan Hajič, Christopher D. Manning, Sampo Pyysalo, Sebastian Schuster, Francis Tyers, and Daniel Zeman. 2020. Universal dependencies v2: An evergrowing multilingual treebank collection. In *Proceedings of the 12th Language Resources and Evaluation Conference (LREC)*.
- Gerard J. Steen, Aletta G. Dorst, J. Berenike Herrmann, Anna A. Kaal, Tina Krennmayr, and Trijntje Pasma. 2010. *A Method for Linguistic Metaphor Identification: From MIP to MIPVU*. John Benjamins.
- William Timkey and Marten van Schijndel. 2021. All bark and no bite: Rogue dimensions in transformer language models obscure representational quality. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Alex Warstadt, Aaron Mueller, Leshem Choshen, Ethan Wilcox, Chengxu Zhuang, Juan Ciro, Rafael Mosquera, Bhargavi Paranjabe, Adina Williams, Tal

722 Linzen, and Ryan Cotterell. 2023. Findings of the  
723 BabyLM challenge: Sample-efficient pretraining on  
724 developmentally plausible corpora. In *Proceedings*  
725 *of the BabyLM Challenge at CoNLL*.

## A Validating the Measure

Experiment 1 plants structure in a *language* and measures a trained model’s representation. A discrepancy there is ambiguous between the measure missing something and the model not mirroring the language, and the second is known to happen: the linear condition cannot represent an interaction at all. This appendix removes the ambiguity by planting  $\mu$ ,  $\alpha$ ,  $\beta$  and  $\gamma$  **directly in a representation**, with no model and no training, and passing the result to the same function the language models are measured with. Ground truth now sits at the level the measure operates on.

The three components are drawn at random and then centered so each satisfies the constraints its role implies:  $\alpha$  and  $\beta$  sum to zero over their own index,  $\gamma$  over both. Without that a planted  $\gamma$  would partly be an item or class main effect and the planted shares would be wrong before the measure saw them. Each is rescaled to hit a target share of the total squared length across the grid, under the same weighting the measure uses. Each mean is then observed  $n$  times under isotropic noise, whose scale is set as a multiple of the mean squared length. The grid crosses the planted shares (the same factorial Experiment 1 runs, mapped into share space, plus the three compositions Experiment 2 measures), item counts from 27 to 180, two and four classes, widths from 8 to 2048, 10, 24 and 100 observations per item and class, three noise levels, planted-orthogonal and planted-non-orthogonal effects, and five seeds: 252,000 runs.

At 24 observations per item and class, inside the 10-to-100 range in which the language-model grids sit, the measured interaction matches the planted one across the whole range (Table 3), and is flagged in essentially every run where anything was planted and about one in twenty where nothing was.

| planted | runs | measured                | flagged |
|---------|------|-------------------------|---------|
| 0.000   | 6300 | 0.0000 [0.0000, 0.0094] | 5%      |
| 0.030   | 900  | 0.0323 [0.0199, 0.0418] | 100%    |
| 0.048   | 1800 | 0.0495 [0.0385, 0.0688] | 100%    |
| 0.056   | 1800 | 0.0573 [0.0458, 0.0918] | 100%    |
| 0.059   | 1800 | 0.0607 [0.0495, 0.1041] | 100%    |
| 0.100   | 1800 | 0.1007 [0.0911, 0.1089] | 100%    |
| 0.111   | 900  | 0.1154 [0.1031, 0.1279] | 100%    |
| 0.167   | 1800 | 0.1708 [0.1594, 0.2204] | 100%    |
| 0.190   | 1800 | 0.1937 [0.1825, 0.2736] | 100%    |
| 0.200   | 1800 | 0.2030 [0.1926, 0.2981] | 100%    |
| 0.240   | 900  | 0.2416 [0.2327, 0.2488] | 100%    |
| 0.333   | 900  | 0.3396 [0.3263, 0.3610] | 100%    |

Table 3: Recovery at 24 observations per item and class in 768 or more dimensions, pooled over noise levels, item counts and widths. Brackets are the 2.5th and 97.5th percentiles across runs.

All three effects read essentially zero when nothing is planted, at every sample size, and each is flagged at about the nominal 5% (Table 4): at 24 observations per item and class a true zero class effect reads 0.0000. This calibration is what licenses reading the small metaphor class effect in Section 3.4 as real rather than as the floor: since a true zero reads essentially zero at this sample size, an interval that excludes zero separates even an effect as small as metaphor’s from zero rather than measuring noise.

| effect       | observations | median  | 95th pct | flagged |
|--------------|--------------|---------|----------|---------|
| class effect | 10           | 0.00000 | 0.0081   | 6.1%    |
| class effect | 24           | 0.00000 | 0.0034   | 6.5%    |
| class effect | 100          | 0.00000 | 0.0009   | 6.3%    |
| interaction  | 10           | 0.00009 | 0.0644   | 6.2%    |
| interaction  | 24           | 0.00004 | 0.0309   | 5.9%    |
| interaction  | 100          | 0.00002 | 0.0083   | 4.9%    |
| item effect  | 10           | 0.00007 | 0.0568   | 4.1%    |
| item effect  | 24           | 0.00003 | 0.0255   | 3.6%    |

(continued)

| effect      | observations | median  | 95th pct | flagged |
|-------------|--------------|---------|----------|---------|
| item effect | 100          | 0.00000 | 0.0064   | 5.0%    |

Table 4: What the measure reports when an effect is planted at exactly zero, by observations per item and class, pooled over every other axis of the validation grid.

The rate at which a genuine zero is nonetheless flagged varies with width and effect rather than holding at the nominal 5% (Table 5). The class effect is the most liberal, and most so at the smallest width, where a true zero is flagged 14% of the time; the item effect is the most conservative, falling to 0% at the widest. No setting exceeds ten percent, and pooled across widths the class and interaction are a little more liberal than the item effect. This is the floor the Experiment 1 detections at small width are read against, so a lone flag at  $d = 8$  is weak on its own. The language-model results sit at  $d \geq 768$ , where the rate is mildly liberal at worst, and each rests on the effect clearing zero in all six models, which even an eight-percent per-model rate would not produce by chance.

| width d | item | class | interaction |
|---------|------|-------|-------------|
| 8       | 3%   | 14%   | 7%          |
| 16      | 7%   | 3%    | 5%          |
| 32      | 4%   | 4%    | 10%         |
| 128     | 6%   | 4%    | 5%          |
| 768     | 2%   | 8%    | 3%          |
| 1024    | 3%   | 5%    | 3%          |
| 2048    | 0%   | 8%    | 8%          |

Table 5: How often a component planted at exactly zero is flagged, by width and by effect.

With the item and class effects planted orthogonal, the measured overlap tracks  $r/d$  closely and sits a little below it (Table 6). This is the number Experiment 2’s overlaps are read against: at  $d = 2048$  a genuinely orthogonal pair reads 0.00096, and the values reported in Section 3.4 are 2.6 to 7.0 times that.

| width d | measured                   | analytic chance |
|---------|----------------------------|-----------------|
| 8       | 0.17501 [0.06900, 0.40018] | 0.25000         |
| 16      | 0.10678 [0.03744, 0.20884] | 0.12500         |
| 32      | 0.05609 [0.01864, 0.10454] | 0.06250         |
| 128     | 0.01714 [0.00450, 0.02611] | 0.02344         |
| 768     | 0.00261 [0.00079, 0.00472] | 0.00391         |
| 1024    | 0.00217 [0.00065, 0.00330] | 0.00293         |
| 2048    | 0.00096 [0.00031, 0.00165] | 0.00146         |

Table 6: Measured item-class overlap when the two effects are planted orthogonal, at 24 observations per item and class. *Measured* is the median over the runs at that width with the 2.5th and 97.5th percentiles in brackets. The *chance* column is the analytic level of Appendix E.

The orthogonal case shows the overlap does not cry wolf; the planted-non-orthogonal case shows it does not miss. With a known fraction of the class direction placed inside the item subspace, the measured overlap rises well above the chance level and falls outside its own null in the large majority of runs (Table 7): at  $d = 2048$  a planted-non-orthogonal pair reads 0.0036 against a chance of 0.00049, detected in 89% of runs. So a reported overlap is one the measure can place at chance when the effects are orthogonal and lift above it when the non-orthogonality is real, which is what makes the Experiment 2 overlaps interpretable in both directions.

| width d | measured                | chance  | detected |
|---------|-------------------------|---------|----------|
| 8       | 0.1688 [0.0824, 0.3998] | 0.12500 | 9%       |

(continued)

| width d | measured                | chance  | detected |
|---------|-------------------------|---------|----------|
| 16      | 0.0951 [0.0414, 0.2026] | 0.06250 | 22%      |
| 32      | 0.0593 [0.0225, 0.1119] | 0.03125 | 36%      |
| 128     | 0.0213 [0.0062, 0.0517] | 0.00781 | 63%      |
| 768     | 0.0065 [0.0011, 0.0298] | 0.00260 | 83%      |
| 1024    | 0.0056 [0.0008, 0.0331] | 0.00098 | 82%      |
| 2048    | 0.0036 [0.0004, 0.0290] | 0.00049 | 89%      |

Table 7: Measured item-class overlap when a fraction of the class direction is planted INSIDE the item subspace (non-orthogonal), at 24 observations per item and class. *Measured* is the median over the runs at that width with the 2.5th and 97.5th percentiles in brackets; *chance* is the analytic random-orientation level; *detected* is the fraction of runs whose overlap falls outside that null’s 95% interval.

## B Capacity and Orthogonality

Whether the item and class effects come out orthogonal depends on whether the representation had room for them to be, which is a fact about capacity rather than about the language. The languages are held fixed across these runs; only the hidden width changes.




Figure 5: Item-class overlap against capacity, by activation, over the 7,560 artificial-language runs. Capacity is the rank the structure needs divided by the hidden width. The dashed line is the median upper end of the null for an arbitrarily oriented effect. Both axes are logarithmic.

Two further properties of the language move the composition, and are given here rather than folded into Figure 3 (Table ef{tbl-lang}). Context competes with the interaction for room: the more of the language that is context, the smaller the interaction’s share of the grid. And the balance the language strikes between the item and the class is reflected in the representation, the measured item share rising with the planted one.

## C A Construction That Is Not Same-Token

Every construction in the body keys the item on the surface form, so the token is identical in both classes of the distinction and the embedding layer has nothing to distinguish them with. That property is what licenses reading the depth profile as contextualization. Applying the same measure to a distinction that lacks the property shows what its absence looks like.

| condition | capacity | runs | interaction                    | item-class overlap     | interaction overlap   |
|-----------|----------|------|--------------------------------|------------------------|-----------------------|
| linear    | 0.125    | 405  | 0.008 [0.002, 0.017] (241/405) | 0.0086 (0.4x, 333/405) | —                     |
| linear    | 0.500    | 405  | 0.004 [0.001, 0.019] (189/405) | 0.0137 (0.2x, 398/405) | —                     |
| linear    | 1.000    | 405  | 0.002 [0.000, 0.024] (57/405)  | 0.0289 (0.2x, 348/393) | —                     |
| linear    | 2.000    | 396  | 0.002 [0.000, 0.556] (111/396) | 0.1049 (0.4x, 188/319) | —                     |
| ReLU      | 0.125    | 405  | 0.083 [0.036, 0.186] (405/405) | 0.0092 (0.6x, 225/405) | 0.087 (1.3x, 316/405) |
| ReLU      | 0.500    | 405  | 0.037 [0.003, 0.073] (367/405) | 0.0408 (0.7x, 231/405) | 0.142 (0.7x, 214/332) |
| ReLU      | 1.000    | 405  | 0.008 [0.000, 0.046] (237/405) | 0.4022 (6.9x, 271/394) | 0.230 (1.1x, 40/97)   |
| ReLU      | 2.000    | 315  | 0.003 [0.000, 0.429] (189/315) | 0.5119 (5.1x, 122/226) | 0.582 (2.4x, 40/77)   |

Table 8: Interaction size and the two overlaps by activation and capacity, over the 7,560 artificial-language runs. Capacity is the rank the structure needs divided by the hidden width; each capacity corresponds to one width, from 128 down to 8. *Interaction* is the median share of the representation over the runs in the row, their interquartile range in brackets, and how many have a re-split 95% interval excluding zero. The overlap columns give the median fraction surviving the projection, then in parentheses how many times its own null median that is and how many runs fall outside that null. The linear arm’s interaction stays below the 2% floor at which orientations are reported, so its last column is empty.

| property                | level        | runs | item  | class | interaction |
|-------------------------|--------------|------|-------|-------|-------------|
| context in the language | under 30%    | 720  | 0.451 | 0.482 | 0.027       |
|                         | 30 to 70%    | 572  | 0.494 | 0.460 | 0.016       |
|                         | over 70%     | 745  | 0.505 | 0.459 | 0.010       |
| item against class      | item < class | 694  | 0.330 | 0.605 | 0.020       |
|                         | item = class | 674  | 0.492 | 0.447 | 0.019       |
|                         | item > class | 669  | 0.643 | 0.311 | 0.014       |

Table 9: How two other properties of the artificial language move the composition of the grid, under ReLU, pooled over widths. Rows within a property partition the same runs. Values are medians of the three shares over the runs in the row.

Number and tense provide that distinction. A plural or a past tense is spelled with an extra subword, so *dog* and *dog-s* are different tokens and the distinction is present in the input before any contextualization occurs.

| family | model                         | feature | layer 0 | deepest | ratio |
|--------|-------------------------------|---------|---------|---------|-------|
| BabyLM | opt-babylm-1.3B-20eps-seed964 | Number  | 0.399   | 0.254   | 0.6   |
| BabyLM | opt-babylm-125m-20eps-seed964 | Number  | 0.344   | 0.221   | 0.6   |
| BabyLM | opt-babylm-350m-20eps-seed964 | Number  | 0.373   | 0.219   | 0.6   |
| Pythia | pythia-1.4b                   | Number  | 0.197   | 0.122   | 0.6   |
| Pythia | pythia-160m                   | Number  | 0.168   | 0.141   | 0.8   |
| Pythia | pythia-410m                   | Number  | 0.164   | 0.120   | 0.7   |
| BabyLM | opt-babylm-1.3B-20eps-seed964 | Tense   | 0.392   | 0.153   | 0.4   |
| BabyLM | opt-babylm-125m-20eps-seed964 | Tense   | 0.305   | 0.137   | 0.5   |
| BabyLM | opt-babylm-350m-20eps-seed964 | Tense   | 0.361   | 0.139   | 0.4   |
| Pythia | pythia-1.4b                   | Tense   | 0.185   | 0.082   | 0.4   |
| Pythia | pythia-160m                   | Tense   | 0.147   | 0.085   | 0.6   |
| Pythia | pythia-410m                   | Tense   | 0.154   | 0.083   | 0.5   |

Table 10: Number and tense, which are marked by an affix and are therefore read at a different subword in the two classes. Columns are as in the deepest-layer table, at layer 0 and at the deepest layer.

The profile inverts (Table 10). Where the same-token constructions begin near zero at layer 0 and grow by one to two orders of magnitude, number and tense begin large and do not grow, declining somewhat with depth. This is what a distinction carried by the input rather than built by the network looks like under this measure, and it is the pattern the body’s constructions were designed to avoid.

The lemma-keyed part of speech figures in Section 3.1 read the same way: what differs between the two keyings is the question the measure is asked, not how well it answers.

## D Model Details

| family | model           | params | width | layers | training data            | tokens |
|--------|-----------------|--------|-------|--------|--------------------------|--------|
| BabyLM | opt-babylm-125m | 125M   | 768   | 12     | BabyLM (child-available) | ~2B    |
|        | opt-babylm-350m | 350M   | 1024  | 24     | BabyLM (child-available) | ~2B    |
|        | opt-babylm-1.3B | 1.3B   | 2048  | 24     | BabyLM (child-available) | ~2B    |
| Pythia | pythia-160m     | 160M   | 768   | 12     | the Pile                 | ~300B  |
|        | pythia-410m     | 410M   | 1024  | 24     | the Pile                 | ~300B  |
|        | pythia-1.4b     | 1.4B   | 2048  | 24     | the Pile                 | ~300B  |

Table 11: The six models, paired by width. Training-token counts are approximate; the OPT-BabyLM models train for 20 epochs over a 100M-word corpus, so their effective token exposure is roughly twenty times the corpus size. *Layers* counts the transformer blocks, and the measure is applied to each block’s output plus the embedding layer.

The two families are matched on width rather than on parameter count, which is what the measure is sensitive to: the chance level for overlap is the class subspace rank over the width (Appendix E), so a comparison across models of different width has to be read against different chance levels. Pairing at 768, 1024 and 2048 removes that from the family comparison.

They differ in two ways that matter elsewhere in the paper. The training corpora differ by more than two orders of magnitude in size, which is the manipulation of interest. And they position tokens differently: Pythia uses rotary embeddings, applied inside attention, while OPT learns absolute position embeddings added at the input. Appendix G uses that difference to rule out absolute position as the source of the interaction.

## E The Decomposition

### E.1 The grid




Figure 6: Each word in each class is observed in many contexts, and those observations are averaged within the pairing to give one entry of the grid. Every entry is a vector in  $\mathbb{R}^d$ , and every item appears in every class.

A word used in a given class is observed in many contexts, and the representation differs a little in each. Averaging those observations within the pairing gives one entry of the grid: for items  $i \in \{1, \dots, I\}$  crossed with classes  $c \in \{1, \dots, C\}$ ,  $M[i, c] \in \mathbb{R}^d$  is the mean representation of item  $i$  in class  $c$  (Figure 6). An item observed in only one class is dropped, as is an item with fewer than 10 tokens in either class, so the grid is complete and every entry rests on a comparable number of observations. The variation across contexts that the averaging sets aside is not discarded; it is what the two splits of Section 2.2.2 are built from.




Figure 7: Centering, drawn for a single hidden dimension read across the whole grid. The all-ones direction is where every position takes the same value, so it carries no differences and no variance, and  $\mu$  is the projection onto it. Removing it leaves a hyperplane of dimension  $N - 1$ , and every remaining component is a displacement inside it. The dashed copy is that displacement moved to the origin, where it lies in the hyperplane.

Two spaces are in play. The representations themselves live in  $\mathbb{R}^d$ . The decomposition happens in the grid index space: fix one hidden dimension, read it across all  $IC$  positions of the grid, and it is a vector in  $\mathbb{R}^{IC}$  (Figure 7). Everything below is that operation carried out once per hidden dimension, with the results stacked back up into the vectors  $\alpha_i$ ,  $\beta_c$  and  $\gamma_{ic}$ .

In that space the all-ones direction is the one in which every position of the grid takes the same value. Moving along it adds the same amount everywhere and creates no difference between any two positions, so it is the direction along which nothing distinguishes anything, and  $\mu$  is the projection onto it. Removing it leaves the orthogonal complement, and the three remaining components are displacements inside that complement rather than raw values.

Each component is fixed by two choices, one per axis: along the item axis either average or center, and along the class axis the same. Averaging keeps what is common and discards what makes positions differ along that axis; centering does the reverse. The four combinations are the four components of Equation 2:

|               | varies over items | varies over classes |
|---------------|-------------------|---------------------|
| $\mu$         | no                | no                  |
| $\alpha_i$    | yes               | no                  |
| $\beta_c$     | no                | yes                 |
| $\gamma_{ic}$ | yes               | yes                 |

So  $\alpha_i$  is what holds of an item in both classes,  $\beta_c$  is what holds of a class for every item, and  $\gamma_{ic}$  is what varies along both at once, which is why it cannot be a fact about either one alone. Since averaging

and centering along an axis add back to what was there, the four reconstruct  $M[i, c]$  exactly, and they are mutually orthogonal.

Centering an axis costs one dimension, the all-ones direction along it. The four components therefore occupy 1,  $I-1$ ,  $C-1$  and  $(I-1)(C-1)$  of the grid's  $IC$  positions, which is why two classes leave  $\beta$  a single direction whatever  $I$  is. Nothing is left over as error:  $\gamma$  is not a residual in the sense of a model that failed to fit, but the part of the representation belonging to the pairing of a particular item with a particular class.

Two consequences are used in the body. First,  $\sum_i \alpha_i = 0$ ,  $\sum_c \beta_c = 0$  and  $\sum_i \gamma_{ic} = \sum_c \gamma_{ic} = 0$  follow from the centering rather than being imposed as constraints. Second, an **additive** grid, one of the form  $M[i, c] = A(i) + B(c)$ , has  $\gamma = 0$  *identically*: writing  $\bar{A}$  and  $\bar{B}$  for the averages,  $\mu = \bar{A} + \bar{B}$ ,  $\alpha_i = A(i) - \bar{A}$ ,  $\beta_c = B(c) - \bar{B}$ , and the remainder cancels exactly. This is the algebra behind the linear condition of Experiment 1: its population grid, the pairing means taken over the same context distribution everywhere, is additive and its interaction is zero. The grid it is measured on is the means of its observed sample, which can also hold a trace of each pairing's own contexts (Experiment 1, Results), so the floor the linear arm provides is read off the planted-zero languages, not deduced.

### E.3 Sizes




Figure 8: How many positions of the grid each component occupies.  $\alpha_i$  does not depend on the class, so it occupies all  $C$  positions of its row;  $\beta_c$  does not depend on the item, so it occupies all  $I$  positions of its column;  $\gamma_{ic}$  takes its own value at every position and is never repeated.

Because the three components are mutually orthogonal, their squared lengths add rather than overlapping, and the total squared length across the grid partitions exactly:

$$\sum_{i,c} \|M[i, c] - \mu\|^2 = C \sum_i \|\alpha_i\|^2 + I \sum_c \|\beta_c\|^2 + \sum_{i,c} \|\gamma_{ic}\|^2,$$

$C$  and  $I$  are not a choice of scaling. They are how many positions of the grid each component occupies (Figure 8):  $\alpha_i$  does not depend on the class, so it sits at all  $C$  positions of its row, and  $\beta_c$  does not depend on the item, so it sits at all  $I$  positions of its column. Every one of the  $I$  items is displaced by the same  $\beta_c$ , so the displacement present in the grid is  $I$  copies of it. The interaction is simply summed because it takes its own value at every position and is never repeated. Dividing all three by  $IC$  would report averages per position instead of totals and leave every share unchanged; what the occupancy counts rule out is weighting the three differently from one another.

Sizes are reported as shares of the total. All three terms scale linearly in  $I$ , so the shares do not depend on how many items the grid contains and are comparable across constructions of different size.

The partition above is of the grid, which has already averaged the contexts away, so it leaves open how much of the representation the grid is. Splitting each observation into its pairing mean and its deviation from it, and using that deviations from a mean sum to zero, gives a second exact partition one level up:

$$\sum_x \|x - \bar{x}\|^2 = \underbrace{\sum_{i,c} \sum_{x \in (i,c)} \|x - \hat{M}[i, c]\|^2}_{\text{within a pairing}} + \underbrace{\sum_{i,c} n_{ic} \|\hat{M}[i, c] - \bar{x}\|^2}_{\text{between pairings}}.$$

The raw between fraction overstates the grid, because  $\hat{M}[i, c]$  is itself an average of  $n_{ic}$  observations and carries  $d\sigma^2/n_{ic}$  of context noise that lands in the between term. The bias falls as tokens per item rise, so it differs across constructions and would not be comparable between them. We therefore report the variance-component form: the within mean square estimates the context variance, the between mean square estimates that plus  $n_0$  times the true spread of the pairing means with  $n_0$  the effective group size of the unbalanced design, and the difference gives the spread alone. Planting no structure at all, the raw fraction reads 0.081 at 12 observations per pairing and 0.005 at 200, while the corrected one reads zero at both.

Because the corrected quantities are variance components rather than weighted sums, they compose with the sizes: each size times the grid’s share gives that component’s share of the representation, and the three plus the context term account for it entirely.

#### E.4 The overlaps and their chance level

The sizes say how much of the representation each component accounts for and nothing about direction. Two components can both be large whether or not they occupy the same directions, and separability is the second question: a word’s representation is an item part plus a class part only if the two can be varied independently, which requires them to lie on distinct axes. That question lives in  $\mathbb{R}^d$ , not in the grid index space the sizes were computed in, and unlike the sizes it does not partition, because each overlap is a relation between two particular components rather than a share of one whole.

Each overlap projects one component into the subspace another occupies:

| projected into the subspace of |                              |
|--------------------------------|------------------------------|
| $\alpha$                       | $\beta$                      |
| $\beta$                        | $\alpha$                     |
| $\gamma$                       | $\alpha$ and $\beta$ jointly |

The first two are not redundant with each other. One direction and thirty-odd are very different targets to land in by chance, and the chance level below is what puts them on a comparable footing.

Subspaces are orthonormal bases from the leading right singular vectors, with rank set by the participation ratio  $r = \text{round}((\sum_k \lambda_k)^2 / \sum_k \lambda_k^2)$  of the scatter eigenvalues, capped at the algebraic maxima above. The participation ratio counts directions discounted by how much each carries: it returns 1 when all the spread lies along one direction and  $k$  when it is even across  $k$  of them. With two classes  $\beta$  spans one direction, so  $r_{\text{class}} = 1$  throughout Experiment 2.

Writing  $S$  for the  $d \times r$  basis and  $V$  for the component being projected, the overlap is  $\|VS\|_F^2 / \|V\|_F^2$ .  $V$  is complete, every vector of the component at full dimension; only  $S$  is rank-truncated. So the rank enters through the target alone, which is why it fixes the chance level without altering any measured overlap.

The overlap is a squared cosine. For one vector and one direction,  $\langle v, u \rangle^2 / \|v\|^2 = \cos^2 \theta$  exactly, so the ratio is a length-weighted mean of  $\cos^2 \theta_i$  over the vectors of  $V$ , with  $\theta_i$  the angle between vector  $i$  and the nearest direction in  $S$ . Weighting by squared length keeps items with small, poorly determined effects from moving the number. The square root inverts to an angle, which is the more legible scale near the floor: an overlap of 0.19 is  $64^\circ$  and one of 0.0056 is  $86^\circ$ , nearly perpendicular rather than merely small.

Squaring rather than summing the signed dot products is not cosmetic. Because  $\sum_i \alpha_i = 0$ , the signed version is  $\langle \sum_i \alpha_i, \beta \rangle = 0$  identically, whatever the geometry: items leaning one way along the class axis cancel those leaning the other. Squaring is also what makes the two directions of a component count

equally, which is what the question requires, since projecting the class direction out costs an item the same either way.

The chance level is not zero. For a randomly oriented  $V$ ,

$$\mathbb{E}[\|VS\|_F^2] = \text{tr}(V^\top V)/d \cdot r \implies \mathbb{E}[\text{overlap}] = r/d,$$

whatever  $V$ 's rank or shape. At  $d = 2048$  that is 0.0005, and Appendix A measures it directly at 0.0010. Any overlap has to be read against that number, which moves by two orders of magnitude across the widths used in this paper.

The measure is exactly invariant to per-dimension rescaling, which the standardization absorbs, and to orthogonal rotation of already-standardized data. It is **not** invariant to general linear reparameterization: we do not whiten, because whitening would make a linear representation read separable in any basis and would discard the basis the model actually computes in. Overlap values are therefore basis-relative, while the sizes and their growth with depth do not depend on the choice.

## F Deepest-Layer Measurements

Table 14 gives the deepest-layer values behind the curves of Figure 4, together with the null each overlap is read against.

| distinction      | model       | item-class overlap | interaction overlap | class effect         | interaction          | portable fraction |
|------------------|-------------|--------------------|---------------------|----------------------|----------------------|-------------------|
| part of speech   | BabyLM 125M | 0.0056 (4.4×)      | 0.172 (3.5×)        | 0.082 [0.081, 0.082] | 0.105 [0.104, 0.106] | 0.44              |
|                  | BabyLM 350M | 0.0052 (5.4×)      | 0.169 (4.7×)        | 0.080 [0.080, 0.080] | 0.091 [0.090, 0.093] | 0.47              |
|                  | BabyLM 1.3B | 0.0032 (6.7×)      | 0.096 (4.5×)        | 0.057 [0.057, 0.057] | 0.114 [0.113, 0.116] | 0.33              |
|                  | Pythia 160M | 0.0058 (4.6×)      | 0.230 (4.2×)        | 0.078 [0.078, 0.078] | 0.232 [0.229, 0.234] | 0.25              |
|                  | Pythia 410M | 0.0067 (7.1×)      | 0.239 (6.1×)        | 0.099 [0.098, 0.099] | 0.245 [0.242, 0.248] | 0.29              |
|                  | Pythia 1.4B | 0.0056 (11.6×)     | 0.209 (9.8×)        | 0.086 [0.086, 0.086] | 0.243 [0.240, 0.246] | 0.26              |
| grammatical role | BabyLM 125M | 0.0038 (3.1×)      | 0.121 (3.6×)        | 0.058 [0.058, 0.058] | 0.044 [0.042, 0.046] | 0.57              |
|                  | BabyLM 350M | 0.0038 (4.2×)      | 0.111 (4.7×)        | 0.056 [0.056, 0.057] | 0.041 [0.039, 0.044] | 0.58              |
|                  | BabyLM 1.3B | 0.0025 (5.4×)      | 0.084 (5.7×)        | 0.047 [0.047, 0.047] | 0.069 [0.066, 0.072] | 0.40              |
|                  | Pythia 160M | 0.0049 (3.9×)      | 0.179 (5.3×)        | 0.078 [0.078, 0.078] | 0.092 [0.088, 0.096] | 0.46              |
|                  | Pythia 410M | 0.0047 (5.0×)      | 0.160 (6.6×)        | 0.074 [0.074, 0.075] | 0.087 [0.084, 0.091] | 0.46              |
|                  | Pythia 1.4B | 0.0062 (12.5×)     | 0.141 (11.1×)       | 0.078 [0.078, 0.079] | 0.094 [0.090, 0.098] | 0.45              |
| metaphor         | BabyLM 125M | —                  | 0.223 (3.9×)        | 0.002 [0.002, 0.002] | 0.025 [0.024, 0.026] | 0.06              |
|                  | BabyLM 350M | —                  | 0.225 (5.6×)        | 0.002 [0.001, 0.002] | 0.023 [0.022, 0.024] | 0.06              |
|                  | BabyLM 1.3B | —                  | 0.128 (4.5×)        | 0.001 [0.001, 0.001] | 0.031 [0.030, 0.032] | 0.04              |
|                  | Pythia 160M | —                  | 0.230 (4.2×)        | 0.003 [0.003, 0.003] | 0.057 [0.055, 0.059] | 0.05              |
|                  | Pythia 410M | —                  | 0.267 (7.8×)        | 0.006 [0.006, 0.006] | 0.073 [0.070, 0.075] | 0.08              |
|                  | Pythia 1.4B | —                  | 0.249 (13.2×)       | 0.006 [0.006, 0.006] | 0.078 [0.075, 0.080] | 0.07              |

Table 14: Deepest-layer measurements, one row per model. The two overlaps are the fraction of the item effect lying in the class subspace and the fraction of the interaction lying in the span of the two main effects, each followed in parentheses by how many times its own rotation null's median it is. That null holds the effect fixed and rotates the target subspace at random 200 times. The class effect and the interaction are shares of the grid, with each model's re-split 95% interval in brackets. The portable fraction is the class effect divided by the sum of the class effect and the interaction.

## G Positional Embeddings Are Not the Source

One architectural explanation of the interaction has to be ruled out before it can be read as a fact about lexical representation. OPT-BabyLM adds a *learned absolute position embedding* at the input. The two classes of a linguistic distinction are not positionally interchangeable in a sentence, since subjects precede objects and nouns and verbs occupy different slots, so a representation that varies with position varies with the class. If that is what the decomposition is picking up, the interaction would be an artifact of where words sit rather than of how they are represented.

We re-extracted every OPT-BabyLM model with *decoder.embed\_positions* zeroed and measured it identically. The ablation was verified per model rather than assumed: with the embedding zeroed, the same token at two different offsets must give identical layer-0 states, and it does, in every model (drift 0.0 to machine precision). Pythia has no such module, so there is nothing to zero there.

The right quantity is the *portable fraction* used in Table 14, the class effect’s share of class-relevant structure. If positions were manufacturing the interaction, removing them should leave a larger share of what distinguishes the two classes in a direction shared across words. Absolute sizes cannot be compared here: zeroing the embeddings in a model that was trained with them is a heavy intervention, and it shrinks every component at once, the class effect included.

| model           | positions | class | interaction | portable fraction |
|-----------------|-----------|-------|-------------|-------------------|
| opt-babylm-125m | zeroed    | 0.012 | 0.034       | 0.264             |
| opt-babylm-125m | intact    | 0.082 | 0.105       | 0.439             |
| opt-babylm-350m | zeroed    | 0.008 | 0.023       | 0.253             |
| opt-babylm-350m | intact    | 0.080 | 0.092       | 0.468             |
| opt-babylm-1.3B | zeroed    | 0.009 | 0.035       | 0.198             |
| opt-babylm-1.3B | intact    | 0.057 | 0.115       | 0.333             |

Table 15: Trained OPT-BabyLM at the deepest layer, with absolute position embeddings intact and zeroed. Sizes are shares of the grid and the portable fraction is the class effect divided by the sum of the class effect and the interaction.

Removing the position embeddings makes these representations *less* separable, not more. The portable fraction falls in every model, from 0.439 to 0.264, from 0.468 to 0.253, and from 0.333 to 0.198 (Table 15). Whatever the absolute position embedding contributes, it is contributing to the part of class structure that *is* carried by a shared direction. Take it away and what remains is more word-specific, not less. The interaction is not an artifact of position.

A second line of evidence points the same way and needs no ablation at all. Pythia never had an absolute position embedding to remove, and its portable fractions are 0.253, 0.287 and 0.262, at or below every intact OPT-BabyLM value. The family without the mechanism shows the phenomenon at least as strongly as the family with it. An explanation resting on absolute position embeddings would have to account for the models that do not have them.

We tested this for part of speech, the construction where the positional worry is sharpest because noun and verb uses of a form occupy systematically different sentence positions. Grammatical role carries the same concern for the same reason, and behaves the same way. Metaphor does not, since literal and figurative uses of a word are not positionally distinguished.